

P2P Systems: Gossip Protocols


CS 6410

By Alane Suhr & Danny Adams

Outline

- ❖ Timeline
- ❖ CAP Theorem
- ❖ Epidemic algorithms for replicated database maintenance
- ❖ Managing update conflicts in Bayou, a weakly connected replicated storage system
- ❖ Conclusion

Timeline



Lamport	Lamport	FLP	Demers	Schneider	Terry	Lamport
Time, Clocks, and the Ordering of Events in a Distributed System	The Byzantine Generals Problem	Impossibility of Distributed Consensus with One Faulty Process	Epidemic algorithms for replicated database maintenance	Implementing fault-tolerant services using the state machine approach: A tutorial	Managing update conflicts in Bayou, a weakly connected replicated storage system	The part-time parliament

CAP

- Consistency -- all nodes contain the same state
- Availability -- requests are responded to *promptly*
- Partition
 - part of a system completely independent from the rest of the system
 - ideally should maintain itself autonomously
- Partition tolerance -- system can stay online and functional even when message passing fails

CAP Theorem Paxos & Gossip


- Paxos: prioritize consistency given a network partition
 - Gossip: prioritize availability given a network partition
-

Gossip



Gossip Overview

- ❑ Authors
- ❑ Motivations
- ❑ Epidemic Models
 - ❑ Direct Mail
 - ❑ Anti-Entropy
 - ❑ Rumor mongering
- ❑ Evaluation
- ❑ DC's
- ❑ Spatial Distribution



A u t h o r s



Alan Demers
Cornell
University



Dan Greene
PARC
Research



Carl Hauser
PhD Cornell
Washington
State University



Scott Shenker
EECS
Berkeley



Doug Terry
Amazon Web
Services

MOTIVATIONS

- Unreliable network
- Unreliable nodes
- CAP: *AP*
 - always be able to respond to a (read/write) request
 - eventual consistency



Epidemic Models




Proposers and Acceptors

- Proposer
 - In Paxos: clients propose an update to the database
 - Epidemic model: a node infects its neighbors
- Acceptor
 - In Paxos: acceptor accepts an update based on one or more proposals
 - Epidemic model: a node is infected by a neighbor


Types of Epidemics

- ❖ Direct Mail
- ❖ Anti-Entropy
- ❖ Rumor Mongering



Advantages

- Simple algorithms
- High Availability
- Fault Tolerant
- Tunable
- Scalable
- Works in Partition




- Notify all neighbors of an update
- Timely and reasonably efficient
- n messages per update




Direct Mail

Direct Mail



Direct Mail



Direct Mail

Messages sent: $O(n)$ where n is number of neighbors


Not fault tolerant -- doesn't guarantee eventual consistency

High volume of traffic with site at the epicenter




Anti-Entropy


- ❑ Site chooses random partner to share data
- ❑ Number of rounds til consistency: $O(\log n)$
- ❑ Sites use custom protocols to resolve conflicts
- ❑ Fault tolerant




Anti-Entropy




Anti-Entropy




Anti-Entropy




Anti-Entropy




Anti-Entropy




Anti-Entropy




Anti-Entropy



Anti-Entropy



Anti-Entropy





What
happens
next?

Mechanism: Push & Pull



Push vs. Pull



What is Push-Pull?



{A, B}

{A, C}

{A, B, C}

{A,B,C}

D

$H(A), H(B)$

$C, H(B)$

B

Propagation times of Push vs. Pull

$$\text{Push: } P_{i+1} = P_i e^{-1}$$

$$\text{Pull: } P_{i+1} = P_i^2$$

Pull is faster!!

P = Probability node hasn't received update after the i^{th} round




Rumor Mongering

1. Sites choose a random neighbor to share information with
2. Transmission rate is tuneable
3. **How long new updates are interesting is also tuneable**
4. Can use push or pull mechanisms



Rumor Mongering Complexity

- $O(\ln n)$ rounds leads to consistency *with high probability*
- Push requires $O(n \ln n)$ transmissions until consistency
- Further proved lower bound for all push-pull transmissions: $O(n \ln \ln n)$




Analogy to epidemiology


- **Susceptible:** site does not know an update yet
- **Infective:** actively sharing an update
- **Removed:** updated and no longer sharing

Rumor mongering: nodes go from *susceptible* to *infective* and eventually (probabilistically) to *removed*


Rumor mongering




Rumor mongering




Rumor mongering




Rumor mongering




Rumor mongering



Rumor mongering



Rumor mongering



Rumor mongering

Pros:

- Fast
- Low call on resources
- Fault-Tolerant
- Less traffic

Cons:

- A site can potentially miss an update



Backups

- Anti-entropy can be used to “update” the network regularly after direct mail or rumor mongering
- If inconsistency found in anti-entropy, run the original algorithm again



Death Certificates

- ❖ How are items deleted using epidemic models?





System Update




Death Certificates

- ❖ How to remove items from epidemic model?
- ❖ Drawbacks
 - Space
 - Increases traffic
 - DC Can be lost
- ❖ Dormant death certificates & retention




Evaluating Epidemic Models

- **Residue:** remaining susceptibles when epidemic finishes
- **Traffic:** $\frac{\text{update traffic}}{\text{number of sites}}$
- **Delay:**
 - T_{avg} : Average time between start of outbreak and arrival of update @ given site
 - T_{last} : Delay until last update



Spatial Distribution

Helping
Or
Hurting




Convergence Times and Traffic

- Linear network: anti entropy
 - Nearest-neighbors
 - $O(n)$ convergence
 - $O(1)$ traffic
 - Random connections
 - $O(\log(n))$ convergence
 - $O(n)$ traffic




Optimizations for realistic network distributions

- Select connections from list of neighbors sorted by distance
- Treat network as linear
- Compute probabilities based on position in list



Rumor Mongering Non-Standard Distribution

- Increase k --
number of rounds
a rumor is
“interesting”
- Use push-pull



Takeaways

- Availability >> consistency
- Updates can be expensive
- Distribution protocols should be robust
- Network design can hurt overall performance
- Byzantine Behavior not addressed

Questions?



Additional Reading

Managing update conflicts in
Bayou, a weakly connected
replicated storage system


1995





- Weak consistency makes unstable network applications possible
- Developing good interfaces allows for complex functions like merging to be interchangeable via the application

Timeline




Lamport	Lamport	FLP	Demers	Schneider	Terry	Lamport
Time, Clocks, and the Ordering of Events in a Distributed System	The Byzantine Generals Problem	Impossibility of Distributed Consensus with One Faulty Process	Epidemic algorithms for replicated database maintenance	Implementing fault-tolerant services using the state machine approach: A tutorial	Managing update conflicts in Bayou, a weakly connected replicated storage system	The part-time parliament

What is Bayou?

- Storage system designed for mobile computing
 - Network is not stable
 - Parts of the network may not be connected all the time
 - **Goal: high availability**
 - Guarantees *weak* consistency






Bayou System Diagram

Consistent Replicas

- Writes are first tentative
- Eventually they are committed, ordered by time
- Clients can tell whether writes are stable (committed)
- Primary servers deal with committing updates




Detecting and Resolving Conflicts

- Dependency checks
- Merge procedures
- Described by the clients,
application-dependent



Conclusions

- Distributed systems need a form of consensus
- Effectively choosing the correct consensus model for a system has to be weighed carefully with the attributes of the system



Acknowledgements

Content Inspired by:

Ki Suh Lee: “Epidemic Techniques”[2009]

Eugene Bagdasaryan: “P2P Gossip Protocols” [2016]

Photos

www.pixabay.com

www.unsplash.com

www.1001freedownloads.com/free-cliparts